

Lecture 12: Bringing it All Together — From Circuits to Execution

(Compilation, Linking, Loading, GOT, and PLT)

Xin Liu

Florida State University

xl24j@fsu.edu

COP 4610 Operating Systems

<https://xinliulab.github.io/cop4610.html>

September 17, 2025

Hardware: Digital Logic Circuits

- The foundation of the digital system world is built on a single axiom: **NAND gates**.
- Using NAND gates, we can:
 - Construct all other logic gates.
 - Build flip-flops to form memory elements.
 - Create any digital logic circuit, including complex systems like mobile applications.

From Simple to Complex:

- Begin with simple digital circuits.
- Gradually scale up to more abstract systems
 - logic gates → circuits → machine code → high-level programming languages.



A	B	Q
0	0	1
0	1	1
1	0	1
1	1	0

Programming as a State Machine

Key Insight:

- All programs (e.g., C, Java, Python) running on a Von Neumann machine are inherently **state machines**.
- The program state consists of:
 - Local variables, global variables, stack, heap, and the program counter (PC).
- Function calls and returns manipulate the state machine:
 - Call: Pushes a new stack frame.
 - Return: Removes the top stack frame.

Debugging Perspective:

- Understanding programs as state machines provides a clear path for debugging:
 - Identify incorrect state transitions (e.g., unexpected values or infinite loops).
 - Diagnose bugs systematically using the state machine model.

How High-Level Code Maps to Hardware:

- Programs like `printf("Hello, World!")` are written as text files (byte sequences in ASCII encoding).
- These byte sequences are:
 - Compiled into machine code instructions.
 - Interpreted and executed by circuits formed with NAND gates.
 - High-level code → Assembly → Machine instructions → Logic gates.

Programming vs. Digital Logic:

- Programming introduces abstraction to interact with the underlying digital logic.
- Tools like compilers and operating systems bridge this gap efficiently.

Instruction Set Architecture

- “Low-level semantics”
- Instructions are based on state machines implemented with logic gates.
- Reference: The RISC-V Instruction Set Manual
 - Easy to implement in hardware.
 - Supports programmable execution effectively.

Compiler (It's a Program Too)

- Translates “high-level state machines”(programs) into “low-level state machines”(instruction sequences).
- Translation principle: Ensures external observable behavior equivalence.

Operating System = Objects + APIs

Application Layer → Libraries → System Calls → Operating System →
Device Drivers → Hardware Abstraction Layer → Physical Hardware →
Logic Gates

All problems in computer science can be solved by another level of indirection, except for the problem of too many layers of indirection. (David Wheeler)

Implementation: Everything is a State Machine

- A state machine operates according to predetermined logic.
- The key lies in defining and initializing the initial state.

Executable Files:

- Define the initial state of a state machine.
- Serve as a **data structure** for process execution.

Key Questions Addressed:

- Q1: How does the operating system load executable files?
- Q2: What is dynamic linking/loading, and how does it work?

Core Topics for This Session

Session Objectives:

- Understand the ELF loader for static binaries.
- Explore dynamic linking and loading mechanisms.

Executable Linkable File (elf)

Making the Program Recognizable to the Machine

Key Manuals for This Lesson:

- **System V ABI:** Defines the System V Application Binary Interface for the AMD64 architecture, providing essential specifications for binary compatibility.
 - [System V ABI \(AMD64 Architecture Processor Supplement\)](#)
- **Refspecs:** Additional reference specifications to deepen understanding of Linux-based systems.
 - [Linux Refspecs](#)

Executable Files: Describing the State Machine

The Operating System: An Execution Environment for Programs (State Machines)

- **Executable File (State Machine Description):**
 - The executable file is a key OS object, describing the initial state and transitions of the program's state machine.
- **Registers:**
 - Most registers are set according to the ABI (Application Binary Interface), with initial setup handled by the OS.
 - For example, the OS initializes the program counter (PC) to start execution.
- **Address Space:**
 - Defined by the binary file and ABI specifications.
 - Includes initial data like `argv` and `envp` (environment variables), along with other necessary information.
- **Additional Information:**
 - The OS may store extra data to aid in debugging and for core dumps in case of errors.

Example: Executable Files on an Operating System

Requirements for an Executable File:

- Must have execution ('x') permission.
- Must be in a format that the loader can recognize as executable.

Example Commands and Output:

```
$ ./a.c
bash: ./a.c: Permission denied

$ ./a.c
bash: ./a.c: Permission denied

$ chmod -x a.out && ./a.out
bash: The file './a.out' is not executable by this user

$ chmod +x a.c && ./a.c
Failed to execute process './a.c'. Reason:
exec: Exec format error
The file './a.c' is marked as an executable but could not be run by
the operating system.
```

Who Decides If a File is Executable?

The Operating System (OS Code - `execve`) Determines Executability:

- The OS, through `execve`, decides whether a file can be executed.

Try It Out:

- Use `strace` to trace `execve` calls and observe execution failures.
 - `strace ./a.c`
 - Without execute permission on `a.c`: `execve` returns `-1`, `EACCES`
 - With execute permission but incorrect format on `a.c`: `execve` returns `-1`, `ENOEXEC`

She-bang (`#!/path/to/interpreter`):

- The She-bang (`# !`) allows specifying an interpreter for a script or executable.
- She-bang effectively performs a “parameter swap” in `execve`, launching the specified interpreter to execute the file.

Example: Running Python Code in a C File

- Save the Following Code as `helloworld.c`:

```
#!/usr/bin/python3
print("Hello_World!")
```

- Give the file execute permission:

```
$ chmod +x helloworld.c
```

- Now, you can directly run the `helloworld.c` file to execute the Python code:

```
$ ./helloworld.c
Hello World!
```

Analyzing Executable Files

GNU Binutils: Essential Tools for Executable Files

- **Creating Executable Files:**

- `ld` (Linker): Combines object files into a single executable.
- `as` (Assembler): Translates assembly code into machine code.
- `ar` and `ranlib`: Manage static libraries.

- **Analyzing Executable Files:**

- `objcopy`, `objdump`, `readelf`: Inspect and modify executables, often used in computer systems basics.
- `addr2line`: Maps addresses to line numbers for debugging.
- `size`, `nm`: Display size information and symbol tables.

Learn More: [GNU Binutils Official Page](#)

Why Can We See All This Information?

Debugging Information Added During Compilation:

- When we compile with debug flags, the compiler includes extra information in the binary.
- This information allows tools like `objdump` and `addr2line` to map assembly code back to the original source code.

Example Command:

- Using `gcc -g -S hello.c` generates assembly code with debugging information.
- This enables us to see additional sections in the assembly output, including variable names, line numbers, and other metadata.

Standard of Debugging Information

Mapping Machine State to “C World” State:

- The DWARF Debugging Standard (dwarfstd.org) defines an instruction set, `DW_OP_XXX`, that is Turing Complete.
- This instruction set can perform “arbitrary computations” to map the current machine state back to the C language state.

Challenges and Limitations:

- **Limited Support for Modern Languages:** Advanced features (e.g., C++ templates) are not fully supported.
- **Complexity of Programming Languages:** As languages evolve, it becomes increasingly challenging to accurately map machine states to source code.
- **Compiler Limitations:** Compilers may not always produce perfect debug information, leading to issues like:
 - Frustrating instances of variables being `<optimized out>`
 - Incorrect or incomplete debugging information

Compilation and Linking

Compilation Process: Assembly and Binary

From High-Level Code to Machine Instructions:

- The compiler generates assembly code (text format).
- The assembler converts it into binary instructions.

Example of Generated Assembly and Binary Code:

```
0000000000000000 <main>:  
0: f3 0f 1e fa endbr64  
4: 48 83 ec 08 sub $0x8,%rsp  
8: 31 c0 xor %eax,%eax  
a: e8 00 00 00 00 callq ????????  
f: 31 c0 xor %eax,%eax  
11: 48 83 c4 08 add $0x8,%rsp  
15: c3 retq
```

When Address Translation is Unavailable:

- For example, during calls to functions like `hello`.
- Temporary placeholders (e.g., `0x0`) are used until the linker resolves the actual address.

Relocation Logic:

- The 4-byte offset must be filled to satisfy the assertion:

```
assert(  
    (char *)hello ==  
    (char *)main + 0xf + // call hello's next PC  
    *(int32_t *)((uintptr_t)main + 0xb) // offset in call  
);
```

Requirements:

- This offset must be written to the executable file.

Relocation Entry in ELF Files

Example Relocation Entry:

Offset	Type	Symbol + Addend
0000000000000000b	R_X86_64_PLT32	hello - 4

Key Points:

- Relocate a 32-bit value to $S + A - P$, where:
 - S: Address of the symbol (e.g., `hello`).
 - A: Addend specified in the relocation entry.
 - P: Address of the relocation itself (`main + 0xb`).
- Understand the behavior of `call S + A - P`.

Understanding Compilation

Compiler (gcc):

- Translates **high-level semantics** (C state machine) into **low-level semantics** (assembly code).

Assembler (as):

- Converts low-level semantics into **binary semantics** (state machine containers).
- Performs a 1-to-1 translation into binary machine code:
 - Includes sections, symbols, and debug information.
- Leaves unresolved information for later (e.g., relocations).

Linker (ld):

- Combines all containers to produce a **complete state machine**.
- Uses:
 - `ldscript (-Wl, --verbose)`
 - Links with C Runtime Objects (CRT).
- Resolves errors related to:
 - Missing symbols.
 - Duplicate symbols.

Static ELF Loader

Implementing ELF Loader on an Operating System

Executable Files

- A data structure describing the initial state (transition) of a state machine.
 - Different from in-memory data structures; "pointers" are replaced by "offsets."
 - Defined in parts within `/usr/include/elf.h`.

Loader

- Parses the data structure, copies it into memory, and jumps to the entry point.
- Creates the initial runtime state of the process (`argv`, `envp`, ...).
 - Example: `loader-static.c`
 - Can load any statically linked code such as `minimal-hello.S`.
 - Correctly processes arguments/environment variables.
- **Reference:** [System V ABI](#) Figure 3.9 (Initial Process Stack).

Loading the Operating System Kernel?

- It is also an ELF file.
- The process includes:
 - Parsing the data structure.
 - Copying it into memory.
 - Jumping to the entry point.

Example Code: `bootmain.c` (Compatible with i386/x86-64 architectures)

Linux Kernel: The Grand Entrance

`bootmain.c` and Linux: Any Fundamental Difference? **None!**

- **Steps to Build the Kernel:**

- Decompress the source package.
- Run `make menuconfig` to generate the `.config` file.
- Compile with `make bzImage -j8`.
- Optionally, patch the kernel (e.g., `kernel/exit.c`).

- **Compilation Results:**

- `vmlinux`: ELF-formatted kernel binary.
 - `vmlinuz`: Compressed image for direct QEMU loading.
 - Use `readelf` to view the entry address at `0x1000000` (physical memory 16MB position).
 - `__startup_64`: RTFSC! Start debugging.
- Always remind yourself: **"Don't panic; it's just a state machine"** (just like your lab work!).

Executable File:

- Describes the initial state of a state machine (**program execution**).
- Unlike in-memory data structures:
 - "Pointers" are replaced by **offsets**.
- Data structure definitions available at:
 - `/usr/include/elf.h`

Loader:

- Parses the data structure, copies it into memory, and performs **jumps**.
- Initializes the process runtime environment:
 - `argv`, `envp`, ...

Reference: [System V ABI](#) Figure 3.9 (Initial Process Stack).

Dynamic Linking and Loading

The Need for "Decoupling Applications"

With increasing library functions, the goal is to enable "runtime linking."

- **Reducing Redundancy:**

- Avoid duplication of library functions in every executable file.
- Ensures efficient disk and memory usage.

- **Versioning Requirements:**

- Follow fundamental conventions, especially respecting function versions.
- Employ ["Semantic Versioning"](#) to manage compatibility.
- Without [version control](#), any new version requires recompiling the **entire program**.

- **Large Project Modularization:**

- Compile parts of a project without relinking the entire program.
- Example: Shared libraries like `libjvm.so`, `libart.so`.
- Analogy: In NEMU, "plugging the CPU into the board."

Challenge of Explaining ELF

- Every year: Hard to explain ELF clearly as it becomes overwhelming.
- Root issue: Closely related concepts are forcibly “dispersed” in data structures.
 - Example: GOT[0], GOT[1], ...

A Simpler Approach

- If the compiler, linker, and loader are under your control:
 - How would you design the most “intuitive” dynamic linking format?
 - Reflect on improvements to reach ELF!
- Assume the compiler generates **position-independent code (PIC)** for you.

Designing a New Binary File Format

Key Idea: Simplify Dynamic Linking with Symbol Tables

- Streamline dynamic linking by focusing on symbol lookup.

Example Structure:

- **DL_HEAD**
 - `LOAD("libc.dl")` # Load dynamic library
 - `IMPORT(putchar)` # Import external symbols
 - `EXPORT(hello)` # Export symbols for dynamic library
- **DL_CODE**
 - Define the function `hello`:
 - `call DSYM(putchar)` # Dynamic linking to symbol
- **DL_END**

Creating a Toolchain for .dl Files with Minimal Effort

Compiler:

- Start with existing tools:
 - Use GCC, GNU `as`.

Binutils:

- Leverage existing tools for basic functionality:
 - `ld` = `objcopy` (borrowed functionality).
 - `as` = GNU `as` (also borrowed).
- Additional tools to customize:
 - `readdl` (like `readelf`).
 - `objdump`.
 - You can adapt tools like `addr2line`, `nm`, `objcopy`, etc.

Critical Component: The Loader

- The loader is essential and needs to be built manually.

Dynamic Linking: Implementation

Header Files:

- `dl.h`: Defines the data structures used for dynamic linking.

Toolchain ("All-in-One" Toolkit):

- `ldbox.c`:
 - Includes tools such as `gcc`, `readelf`, `objdump`, and `interp`.

Example Code:

- `libc.S`: Provides implementations for `putchar` and `exit`.
- `libhello.S`: Calls `putchar` and provides `hello`.
- `main.S`: Calls `hello` and provides `main`.
- (Assumes that your high-level language compiler can generate assembly code in this format.)

Return to ELF

Addressing Design Flaws in .dl Files

Memory Protection and Load Address:

- Allow mapping parts of the .dl file to specific memory locations with controlled permissions (via program header table).

Custom Loader Specification:

- Permit the use of custom loaders instead of relying solely on `dlbox`.
- Introduce `INTERP`.

Space Efficiency:

- Store strings in a constant pool and access them through a unified "pointer" mechanism.
- This contributes to the complexity of reading ELF files.

Additional Improvements:

- Less critical but still valuable enhancements.
- Follow documentation: `RTFM/RTFSC`.

DSYM as Indirect Memory Access

```
#define DSYM(sym) *sym(%rip)

extern void foo();
foo();
```

One Syntax, Two Scenarios:

- **From Other Compilation Units (Static Linking):**
 - Direct PC-relative jump.
- **Dynamic Linking Libraries:**
 - Symbol resolution is mandatory (cannot be determined at compile-time).

The "Invention" of GOT & PLT

Our "Symbol Table" is the Global Offset Table (GOT)

- Now you won't misunderstand the concept of GOT!
- The concept and name aren't important; the process of invention is.

Unifying Static and Dynamic Linking: Both Use Static!

- Add a layer of indirection: Procedure Linkage Table (PLT).
- All unresolved symbols are unified and translated into `call`.
- Modern processors optimize this kind of jump (e.g., Branch Target Buffer (BTB)).

Example:

```
putchar@PLT:
    call DSYM(putchar) # in ELF: jmp *GOT[n]

main:
    call putchar@PLT
```

Revisiting `printf`

- You will find that our “minimal” binary file format is almost exactly the same!
- ELF even includes some additional hacks (e.g., lazy binding).

```
00000000000010c0 <printf@plt>:  
10c0: endbr64  
10c4: bnd jmpq *0x2efd(%rip) # DSYM(printf)  
10cb: nopl 0x0(%rax,%rax,1)  
  
00000000000011c9 <main>:  
...  
1246: callq 10c0 <printf@plt>
```


One Last Question: Data

- What if we want to reference data in dynamically linked libraries?
 - Data cannot add another layer of indirection
- Issues with `stdout/errno/envIRON`:
 - Multiple libraries use them, but there should only be **one copy**!

Of course, we did an experiment!

- Use `readelf` to check if `stdout` differs
- Set a watch point in `gdb`
 - It turns out to be treated specially
 - It's a kind of “workaround”

Summary: What is an Executable File?

Before Learning Operating Systems:

- "Something you double-click to open a window."



After Learning Operating Systems:

- An **object** in the operating system (a file).
- A **sequence of bytes** (which can be interpreted as text or encoded symbols).
- A data structure describing the **initial state of a state machine**.

Questions Answered in This Lecture:

- **Q1:** How are executable files loaded by the operating system?
- **Q2:** What are dynamic linking and dynamic loading?

Take-away Messages:

- **Loader:**
 - Uses the underlying mechanism to "transport" data structures according to specifications
- **Dynamic Linking/Loading:**
 - GOT, PLT, and the smallest binary format