

# Lecture 24: Device Driver


## (Abstraction, Design, CUDA)

Xin Liu

Florida State University  
xliu15@fsu.edu

COP 4610 Operating Systems  
<https://xinliulab.github.io/FSU-COP4610-Operating-Systems/>

- From the OS' perspective, an I/O device is essentially  
**A Set of Registers and Protocols**
- Regardless of what is connected behind them
  - Examples: serial port, keyboard, printer



- **Connecting Devices to the Computer:**
  - **Bus:** The sole external interface of the computer, connecting all devices.
    - The bus itself can be viewed as a device.
- **Managing the Bus:**
  - **Interrupt Controller:**
    - Coordinates device interrupts to manage CPU attention.
    - Avoids busy-waiting
  - **DMA:** Acts as a co-processor, assisting in data transfer for I/O devices.
- **Expanding Co-Processing:**
  - **GPU:** Another specialized co-processor.
  - Enables **heterogeneous computing** alongside the CPU.

- Each type of device has its own protocol and set of registers.
- Even for the same type of device, different models may have different registers and protocols.
- Directly exposing devices to the OS greatly increases the likelihood of errors.

## Solution: Device Abstraction

To address this, we abstract all devices to communicate with the CPU in a unified manner as much as possible.

Today's Key Question:

How does the OS enable applications to access these devices?

- **What is a device driver?**
- **How do we abstract the device?**

Main Topics for Today:

- **Principles of Device Drivers**
- **Linux Device Driver Design**
- **GPU and CUDA**

# Principles of Device Drivers

## Input and Output

- “Readable (read) and writable (write) byte sequences (streams or arrays)”
- Most common devices fit this model:
  - Character Device:
    - Terminal/Serial Port - Byte Stream
    - Printer - Byte Stream (e.g., PostScript files)
  - Block Device:
    - Hard Disk - Byte Array (Block Access)
  - GPU: Neither a Character Device Nor a Block Device
    - Byte Stream (Control) + Byte Array (Display)

# How to Abstract I/O Devices

Since I/O devices are designed for input and output, we can abstract them using basic operations:

- Devices as Objects (Files) Supporting Various Operations
  - System call **read** - Reads data from a specified location on the device
  - System call **write** - Writes data to a specified location on the device
  - System call **ioctl** - Reads/Sets the status of the device
    - RTFM: `man 2 ioctl`

# What is a Device Driver?

Translates system calls (read/write/ioctl/...) into interactions with device registers

- Essentially a piece of kernel code
- Human → Shell → System Calls → Driver → Devices
  - Our computer continuously performs abstraction and translation
- May block (e.g., waiting on a semaphore P operation, awakened by V operation in interrupt)

# Examples of Objects in /dev/

```
$ cd /dev && ls
```

- **Note:** A driver is simply code:
  - It may not have a real device behind it.
  - It can simulate devices entirely through code.
- Examples:
  - /dev/random, /dev/urandom - Random number generators
    - Try: head -c 512 /dev/urandom | xxd
    - May use hardware device for true random numbers, or software for pseudo-random numbers
  - /dev/null - "Null" device
    - Try: yes > /dev/null (Fake write success)
    - Try: cat /dev/null (No data to read)
    - Observe using strace
  - /dev/zero - "Zero" device
    - Try: cat /dev/zero | head -c 512 | xxd

## I/O devices appear as a "black box"

- Prone to Errors!
- Any code mistake simply causes it to "not work"
- Device drivers: often the **lowest** quality code in Linux kernel

## The Complexity of Device Drivers

- The computer industry includes countless devices, each with unique registers and protocols.
- Often, only the device's programmers understand them, and even they may not fully.
- Since Vista, Microsoft, and now Linux, have been moving drivers to user space to prevent system crashes from driver errors.

# Limitations of Byte Stream Abstraction

Devices involve not only data but also **control**

- Especially for additional functions and configurations of devices
- All extra features rely on `ioctl`
  - Arguments, returns, and semantics of `ioctl()` vary according to the device driver in question
  - Complex “hidden specifications”

## Examples

- Printer settings: print quality, paper feed, duplex, card stock, cleaning, auto binding, ...
  - A printer worth tens of thousands is not that simple !
- Keyboard lighting effects, repeat rate, macro programming, ...
- Disk health, cache control, ...

## Functions beyond the "Byte Stream"

- Devices are far more complex than simple input/output streams.
- Many features require specific configurations, often handled through `ioctl` calls.
- This complexity is just the tip of the iceberg.

# The Terminal: A Special Device in UNIX

- The Shell acts as the interface between the user and the OS.
- The **terminal** is the I/O device behind this interface.
- Terminals support text-based input and output, enabling users to interact with the OS.
- They handle complex features like:
  - Input modes, signal handling, and line editing
  - Terminal settings, which can be configured dynamically
    - **RTFM:** `tty`, `stty`, ...



# The Terminal Usage

```
$ echo hello > /dev/pts/1
```

- Find out which system call “recognizes” the terminal.

```
#include <stdio.h>

int main() {
    printf("Hello, World\n");
}
```

- Tip: Compare using the following two commands:

```
$ strace ./a.out
```

```
$ strace ./a.out > /dev/null
```

# More On Terminal

## Terminal as a Line Buffer

- Terminal uses **line buffering**: output is stored until a newline ('\n') or 'fflush(stdout);' is called.
- Files, on the other hand, use **full buffering**, meaning data is stored until the buffer is full or manually flushed.
- Buffering improves performance by reducing the number of slow system calls.
- If an error occurs before flushing, buffered data may not be displayed.

```
#include <stdio.h>

int main() {
    printf("Hello
World!");

    int *p;
    p = NULL;
    *p = 1;
}
```

```
#include <stdio.h>

int main() {
    printf("Hello
World!\n");

    int *p;
    p = NULL;
    *p = 1;
}
```

```
#include <stdio.h>

int main() {
    printf("Hello
World!");
    fflush(stdout);
    int *p;
    p = NULL;
    *p = 1;
}
```



# Understanding Terminal Buffering

- Consider the following code:

```
#include <stdio.h>
#include <unistd.h>
#include <sys/wait.h>

int main(int argc, char *argv[]) {
    for (int i = 0; i < 2; i++) {
        fork();
        printf("Hello World!\n");
    }
    for (int i = 0; i < 2; i++) {
        wait(NULL);
    }
}
```

- When we run `./a.out`, we get 6 "Hello World!"
- When we run `./a.out | cat`, we get 8 "Hello World!"

# Understanding Fork

- `fork()` duplicates everything in the process, including the output buffers!
- When we run `./a.out | cat`, we get 8 "Hello World!"
  - It's connected to a pipe, not directly to the terminal.
  - The buffer does not flush after each newline; it waits until it's full or the program ends.
  - When `fork()` is called, the parent process's buffer (which may contain "Hello World!") is duplicated in the child process.
  - Both parent and child processes may flush the same buffered data, leading to duplicated outputs.
  - Result: Buffer duplication causes additional outputs; 8 "Hello World!" are printed.

# Principles of Device Driver Design

How can we use a computer to launch a nuclear missile?

Have A Try: [launcher](#)

# 1. Define the Driver's Purpose

## Key Questions to Define:

- What operations will the driver support? (e.g., read, write)
- How will the device interact with the user or system?
- Example: Our nuclear missile code triggers an ASCII art when the password is received.

## 2. Register Device Number and Class

- **Major Number:** Specifies the driver type that handles the device.
  - Tells the OS which driver to use for this type of device.
  - Assigned by `alloc_chrdev_region` and stored in `dev_major`.
- **Minor Number:** Differentiates instances of the same type of device.
  - Allows the driver to handle multiple devices of the same type.
  - Specified by `MKDEV(dev_major, i)`, creating instances like '`/dev/nuke0`' (Minor 0) and '`/dev/nuke1`' (Minor 1).

```
ret = alloc_chrdev_region(&dev, 0, MAX_DEV, "nuke");
if (ret < 0) {
    printk(KERN_ALERT "nuke: Failed to allocate char
device region\n");
    return ret;
}

dev_major = MAJOR(dev);

lx_class = class_create(THIS_MODULE, "nuke");
```

### 3. Define File Operations Interface

#### File Operations:

- Define how users will interact with the device:
  - **read:** Retrieves data from the device.
  - **write:** Sends data to the device.
  - **open/release:** Opens and closes access to the device.

```
static ssize_t lx_read(struct file *, char __user *,  
    size_t, loff_t *);  
static ssize_t lx_write(struct file *, const char __user  
    *, size_t, loff_t *);  
  
static struct file_operations fops = {  
    .owner = THIS_MODULE,  
    .read = lx_read,  
    .write = lx_write,  
};
```

## 4. Implement Core Functions

### Core Functions Based on file\_operations:

- Each function provides specific device behavior:
  - **lx\_read:** Returns "This is dangerous!" when read.
  - **lx\_write:** Checks input for specific pattern to trigger ASCII art.

```
static ssize_t lx_write(struct file *file, const char
__user *buf, size_t count, loff_t *offset) {
.....
if (cmp_result == 0) {
    printk(KERN_INFO "nuke: correct password entered.\n")
;
    const char *EXPLODE[] = {
        //Your ASCII art
    };
    int i;
    for (i = 0; i < sizeof(EXPLODE) / sizeof(EXPLODE[0]);
i++) {
        printk(KERN_INFO "%s\n", EXPLODE[i]);
    }
} else {
    printk(KERN_INFO "nuke: incorrect secret, cannot

```

## 5. Initialization and Cleanup

- `module_init`: Registers device and sets up file operations.
- `module_exit`: Unregisters device and cleans up resources.

```
static void __exit lx_exit(void) {
    device_destroy(lx_class, MKDEV(dev_major, 0));
    class_unregister(lx_class);
    class_destroy(lx_class);
    unregister_chrdev_region(MKDEV(dev_major, 0),
        MINORMASK);
}

module_init(lx_init);
module_exit(lx_exit);
```

## Key Steps in Device Driver Design:

- Define purpose and main functions.
- Register device (major and minor numbers, class).
- Set up file operations interface (read, write, ioctl).
- Implement core functions for device-specific logic.
- Initialize resources and provide cleanup routines.

# Building the Driver

- Use a Makefile to compile into a '.ko' file:

```
obj-m += nuke.o

KDIR := /lib/modules/$(shell uname -r)/build

all:
    make -C $(KDIR) M=$(PWD) EXTRA_CFLAGS=-Wno-error
modules

clean:
    make -C $(KDIR) M=$(PWD) clean
```


# Install the Driver

- make
- Load module: `sudo insmod nuke.ko`
- Check '/dev' for device files ('/dev/nuke0', '/dev/nuke1') using  
`ls -l /dev/nuke*`
- Set the devices to be readable and writable with  
`chmod 666 /dev/nuke0`
- Test with `echo "COP4610" | sudo tee /dev/nuke0`
- Observe logs with `dmesg | tail -n 20`
- Unload the module with `sudo rmmod nuke`

# CUDA: Programming for GPUs

## Single Instruction, Multiple Threads

- Many threads execute the same instruction
- Each thread has some thread-local data (e.g., ID)
- Highly optimized design
  - One Program Counter (PC), multiple data elements
  - Successor to VLIW and SIMD architectures



# Introduction to CUDA Toolchain

**Reference:** Parallel Thread Execution ISA Application Guide

- CUDA uses an instruction set architecture (ISA) for parallel thread execution.
- The code is compiled into SASS (machine code).
  - Use `cuobjdump --dump-ptx` or `--dump-sass` to view assembly code.

## Essential Tools in the CUDA Toolchain

- **gcc** → **nvcc**: NVIDIA CUDA Compiler, used to compile CUDA code.
- **binutils** → **cuobjdump**: Disassembler for CUDA binaries.
- **gdb** → **cuda-gdb**: Debugger for CUDA code.
  - Allows debugging directly on the GPU!
- **perf** → **nvprof**: Performance profiling tool to analyze GPU code.
- ...

CUDA brings a complete set of tools to the GPU, similar to what we have for CPUs.



## GPU Drivers are Complex

- Complete Toolchain
  - Just-in-time (JIT) compilation
  - Profiler
  - ...
- API Implementation
  - `cudaMemcpy`, `cudaMalloc`, ...
  - Kernel execution
  - Mostly implemented via `ioctl`
- Device Compatibility

**Note:** NVIDIA open-sourced its [driver](#) in 2022!

- Before that ... [Video](#)

- What are device drivers?
  - Translate `read/write/ioctl` calls into device-understandable protocols
- Device Driver Design Principles